

Lab 4

Amy Butler

11:59PM March 10, 2021

Load up the famous iris dataset. We are going to do a different prediction problem. Imagine the only input x is Species and you are trying to predict y which is Petal.Length. A reasonable prediction is the average petal length within each Species. Prove that this is the OLS model by fitting an appropriate `lm` and then using the `predict` function to verify.

#We need to prove the predictions will match the \bar{y} of each species #The numbers match so this worked #Find the \bar{y} for each species

```
data(iris)

mod=lm(Petal.Length ~ Species, iris)

mean(iris$Petal.Length[iris$Species == "setosa"])
## [1] 1.462

mean(iris$Petal.Length[iris$Species == "versicolor"])
## [1] 4.26

mean(iris$Petal.Length[iris$Species == "virginica"])
## [1] 5.552

predict(mod, data.frame(Species=c("setosa")))
##      1
## 1.462

predict(mod, data.frame(Species=c("versicolor")))
##      1
## 4.26

predict(mod, data.frame(Species=c("virginica")))
##      1
## 5.552
```

Construct the design matrix for the previous linear model with an intercept, X , without using `model.matrix`.

#A design matrix makes column vectors that measure what we care about (in this case species) and a column of ones to fit an intercept.

```
X=cbind(1,iris$Species=="versicolor",iris$Species=="virginica" )
head(X)
```

```
##      [,1] [,2] [,3]
## [1,]    1    0    0
## [2,]    1    0    0
## [3,]    1    0    0
## [4,]    1    0    0
## [5,]    1    0    0
## [6,]    1    0    0
```

Find the hat matrix H for this regression.

#%*% is matrix multiplication #Solve() finds the inverse of a matrix #t(X) is X transpose
#The rank should be three because there are 3 columns in X and the projectin of yhat onto the column space of X should be a linear combination of the 3 columns of X.

```
H=X %*% solve(t(X) %*% X) %*% t(X)
Matrix::rankMatrix(H)
```

```
## [1] 3
## attr(,"method")
## [1] "tolNorm2"
## attr(,"useGrad")
## [1] FALSE
## attr(,"tol")
## [1] 3.330669e-14
```

Verify this hat matrix is symmetric using the expect_equal function in the package testthat.

#Pacman loads a package #If there is no error then it worked and the matrix is symmetric

```
pacman::p_load(testthat)
expect_equal(H, t(H))
```

Verify this hat matrix is idempotent using the expect_equal function in the package testthat.

#No output means it works and the matrix is idempotent

```
expect_equal(H, H%*%H)
```

Using the diag function, find the trace of the hat matrix.

#Diag returns, extracts, or replaces the diagonal of a matrix #Trace is the sum of the diagonal #The trace is also the rank

```
sum(diag(H))
```

```
## [1] 3
```

It turns out the trace of a hat matrix is the same as its rank! But we don't have time to prove these interesting and useful facts..

For masters students: create a matrix X_{\perp} .

#TO-DO

Using the hat matrix (H), compute the \hat{y} vector and using the projection onto the residual space, compute the e vector and verify they are orthogonal to each other.

#H%*% does a projection #The table should display the 50 \hat{y} 's for each of the three species #e should be a column vector of decimals with no pattern

```
y=iris$Petal.Length
y_hat= H %*% y
table(y_hat)

## y_hat
## 1.462  4.26 5.552
##    50    50    50

e= (diag(nrow(iris))-H) %*% y
e

##           [,1]
## [1,] -0.062
## [2,] -0.062
## [3,] -0.162
## [4,]  0.038
## [5,] -0.062
## [6,]  0.238
## [7,] -0.062
## [8,]  0.038
## [9,] -0.062
## [10,]  0.038
## [11,]  0.038
## [12,]  0.138
## [13,] -0.062
## [14,] -0.362
## [15,] -0.262
## [16,]  0.038
## [17,] -0.162
## [18,] -0.062
## [19,]  0.238
## [20,]  0.038
## [21,]  0.238
## [22,]  0.038
## [23,] -0.462
## [24,]  0.238
## [25,]  0.438
## [26,]  0.138
```

```
## [27,] 0.138
## [28,] 0.038
## [29,] -0.062
## [30,] 0.138
## [31,] 0.138
## [32,] 0.038
## [33,] 0.038
## [34,] -0.062
## [35,] 0.038
## [36,] -0.262
## [37,] -0.162
## [38,] -0.062
## [39,] -0.162
## [40,] 0.038
## [41,] -0.162
## [42,] -0.162
## [43,] -0.162
## [44,] 0.138
## [45,] 0.438
## [46,] -0.062
## [47,] 0.138
## [48,] -0.062
## [49,] 0.038
## [50,] -0.062
## [51,] 0.440
## [52,] 0.240
## [53,] 0.640
## [54,] -0.260
## [55,] 0.340
## [56,] 0.240
## [57,] 0.440
## [58,] -0.960
## [59,] 0.340
## [60,] -0.360
## [61,] -0.760
## [62,] -0.060
## [63,] -0.260
## [64,] 0.440
## [65,] -0.660
## [66,] 0.140
## [67,] 0.240
## [68,] -0.160
## [69,] 0.240
## [70,] -0.360
## [71,] 0.540
## [72,] -0.260
## [73,] 0.640
## [74,] 0.440
## [75,] 0.040
## [76,] 0.140
```

```
## [77,] 0.540
## [78,] 0.740
## [79,] 0.240
## [80,] -0.760
## [81,] -0.460
## [82,] -0.560
## [83,] -0.360
## [84,] 0.840
## [85,] 0.240
## [86,] 0.240
## [87,] 0.440
## [88,] 0.140
## [89,] -0.160
## [90,] -0.260
## [91,] 0.140
## [92,] 0.340
## [93,] -0.260
## [94,] -0.960
## [95,] -0.060
## [96,] -0.060
## [97,] -0.060
## [98,] 0.040
## [99,] -1.260
## [100,] -0.160
## [101,] 0.448
## [102,] -0.452
## [103,] 0.348
## [104,] 0.048
## [105,] 0.248
## [106,] 1.048
## [107,] -1.052
## [108,] 0.748
## [109,] 0.248
## [110,] 0.548
## [111,] -0.452
## [112,] -0.252
## [113,] -0.052
## [114,] -0.552
## [115,] -0.452
## [116,] -0.252
## [117,] -0.052
## [118,] 1.148
## [119,] 1.348
## [120,] -0.552
## [121,] 0.148
## [122,] -0.652
## [123,] 1.148
## [124,] -0.652
## [125,] 0.148
## [126,] 0.448
```

```
## [127,] -0.752
## [128,] -0.652
## [129,]  0.048
## [130,]  0.248
## [131,]  0.548
## [132,]  0.848
## [133,]  0.048
## [134,] -0.452
## [135,]  0.048
## [136,]  0.548
## [137,]  0.048
## [138,] -0.052
## [139,] -0.752
## [140,] -0.152
## [141,]  0.048
## [142,] -0.452
## [143,] -0.452
## [144,]  0.348
## [145,]  0.148
## [146,] -0.352
## [147,] -0.552
## [148,] -0.352
## [149,] -0.152
## [150,] -0.452
```

Compute SST, SSR and SSE and R^2 and then show that $SST = SSR + SSE$.

#The SSE & SST can be written in the previous formula format we used or in matrix form like below #No output of the expect_equals means that SSR+SSE does equal SST

```
y_bar=mean(y)

SSE= t(e) %*% e
SSE

##           [,1]
## [1,] 27.2226

SST= t(y-y_bar) %*% (y-y_bar)
SSR= t(y_hat-y_bar) %*% (y_hat-y_bar)
RSQ= 1-SSE/SST
RSQ

##           [,1]
## [1,] 0.9413717

expect_equal(SSR+SSE, SST)
```

Find the angle θ between $y - \bar{y}1$ and $\hat{y} - \bar{y}1$ and then verify that its cosine squared is the same as the R^2 from the previous problem.

#Theta should be close to zero #pi/180 turns theta into degrees

```
theta = acos(t(y-y_bar) %*% (y_hat-y_bar) / sqrt(SST * SSR))
theta * (180/pi)

##           [,1]
## [1,] 14.01245
```

Project the y vector onto each column of the X matrix and test if the sum of these projections is the same as yhat.

#We want this to fail when adding the projections

```
proj1= (X[,1] %*% t(X[,1]) / as.numeric(t(X[,1]) %*% X[,1])) %*% y
proj2= (X[,2] %*% t(X[,2]) / as.numeric(t(X[,2]) %*% X[,2])) %*% y
proj3= (X[,3] %*% t(X[,3]) / as.numeric(t(X[,3]) %*% X[,3])) %*% y
```

Construct the design matrix without an intercept, X, without using model.matrix.

```
anova_mod = lm(Petal.Length ~ 0 + Species, iris)
```

Find the OLS estimates using this design matrix. It should be the sample averages of the petal lengths within species.

```
b = solve(t(X)%*%X)%*%t(X)%*%y
Model=lm(Petal.Length~X,iris)
Model

##
## Call:
## lm(formula = Petal.Length ~ X, data = iris)
##
## Coefficients:
## (Intercept)          X1          X2          X3
##      1.462         NA       2.798       4.090
```

Verify the hat matrix constructed from this design matrix is the same as the hat matrix constructed from the design matrix with the intercept. (Fact: orthogonal projection matrices are unique).

```
X = cbind(as.integer(iris$Species == "setosa"), as.integer(iris$Species ==
"versicolor"), as.integer(iris$Species == "virginica"))
H_second = X %*% solve(t(X) %*% X) %*% t(X)
```

Project the y vector onto each column of the X matrix and test if the sum of these projections is the same as yhat.

#Same format as problem above with projections

```
proj1 = ((X[,1] %*% t(X[,1])) / as.numeric(t(X[,1]) %*% X[,1])) %*% y
proj2 = ((X[,2] %*% t(X[,2])) / as.numeric(t(X[,2]) %*% X[,2])) %*% y
proj3 = ((X[,3] %*% t(X[,3])) / as.numeric(t(X[,3]) %*% X[,3])) %*% y
```

Convert this design matrix into Q , an orthonormal matrix.

$$\begin{aligned} \text{qrX} &= \text{qr}(X) \\ Q &= \text{qr}.Q(\text{qrX}) \end{aligned}$$

Project the y vector onto each column of the Q matrix and test if the sum of these projections is the same as y_{hat} .

```
proj1 = ((Q[,1] %>% t(Q[,1])) / as.numeric(t(Q[,1] %>% Q[,1])) %>% y
proj2 = ((Q[,2] %>% t(Q[,2])) / as.numeric(t(Q[,2] %>% Q[,2])) %>% y
proj3 = ((Q[,3] %>% t(Q[,3])) / as.numeric(t(Q[,3] %>% Q[,3])) %>% y
```

Find the $p = 3$ linear OLS estimates if Q is used as the design matrix using the `lm` method. Is the OLS solution the same as the OLS solution for X ?

```
lm(Petal.Length~Q[,3],iris)

##
## Call:
## lm(formula = Petal.Length ~ Q[, 3], data = iris)
##
## Coefficients:
## (Intercept)      Q[, 3]
##      2.861      -19.028
```

Use the predict function and ensure that the predicted values are the same for both linear models: the one created with X as its design matrix and the one created with Q as its design matrix.

[illegible]


```
##      66      67      68      69      70      71      72      73      74      75      76      77
78
## 4.260 4.260 4.260 4.260 4.260 4.260 4.260 4.260 4.260 4.260 4.260 4.260
4.260
##      79      80      81      82      83      84      85      86      87      88      89      90
91
## 4.260 4.260 4.260 4.260 4.260 4.260 4.260 4.260 4.260 4.260 4.260 4.260
4.260
##      92      93      94      95      96      97      98      99     100     101     102     103
104
## 4.260 4.260 4.260 4.260 4.260 4.260 4.260 4.260 4.260 5.552 5.552 5.552
5.552
##     105     106     107     108     109     110     111     112     113     114     115     116
117
## 5.552 5.552 5.552 5.552 5.552 5.552 5.552 5.552 5.552 5.552 5.552 5.552
5.552
##     118     119     120     121     122     123     124     125     126     127     128     129
130
## 5.552 5.552 5.552 5.552 5.552 5.552 5.552 5.552 5.552 5.552 5.552 5.552
5.552
##     131     132     133     134     135     136     137     138     139     140     141     142
143
## 5.552 5.552 5.552 5.552 5.552 5.552 5.552 5.552 5.552 5.552 5.552 5.552
5.552
##     144     145     146     147     148     149     150
## 5.552 5.552 5.552 5.552 5.552 5.552 5.552
```

Clear the workspace and load the boston housing data and extract X and y . The dimensions are $n = 506$ and $p = 13$. Create a matrix that is $(p + 1) \times (p + 1)$ full of NA's. Label the columns the same columns as X . Do not label the rows. For the first row, find the OLS estimate of the y regressed on the first column only and put that in the first entry. For the second row, find the OLS estimates of the y regressed on the first and second columns of X only and put them in the first and second entries. For the third row, find the OLS estimates of the y regressed on the first, second and third columns of X only and put them in the first, second and third entries, etc. For the last row, fill it with the full OLS estimates.

```
rm(list=ls())
Boston=MASS::Boston
X=cbind(1, as.matrix(Boston[,1:13]))
y=Boston[,14]
p1=ncol(X)
matrixp1=matrix(NA, nrow=p1, ncol=p1)
for(j in 1:ncol(X)){
  Xj=X[,1:j]
  matrixp1[j,1:j]=solve(t(Xj)%*%Xj)%*%t(Xj)%*%y
}
```

Why are the estimates changing from row to row as you add in more predictors?

Because the predictions are getting better and better with more data.

Create a vector of length $p + 1$ and compute the R^2 values for each of the above models.

```
vector=c(1:14)
for(i in 1:ncol(X)){
  model=lm(y~X[,1:ncol(X)])
  vector[i]=summary(model)$r.squared
}
```

Is R^2 monotonically increasing? Why?

R-squared is monotonically increasing because with more data there will be a better explanation.